

cargar el conjunto de datos de entrenamiento

```
In [ ]: import pandas as pd
```

```
bike_data = pd.read_csv('daily-bike-share.csv')  
bike_data.head()
```

```
Out[ ]:
```

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	aten
0	1	1/1/2011	1	0	1	0	6	0	2	0.344167	0.3636
1	2	1/2/2011	1	0	1	0	0	0	2	0.363478	0.3537
2	3	1/3/2011	1	0	1	0	1	1	1	0.196364	0.1894
3	4	1/4/2011	1	0	1	0	2	1	1	0.200000	0.2121
4	5	1/5/2011	1	0	1	0	3	1	1	0.226957	0.2292

```
In [ ]: bike_data['day'] = pd.DatetimeIndex(bike_data['dteday']).day  
bike_data.head(32)
```

Out[]:

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	at
0	1	1/1/2011	1	0	1	0	6	0	2	0.344167	0.36
1	2	1/2/2011	1	0	1	0	0	0	2	0.363478	0.35
2	3	1/3/2011	1	0	1	0	1	1	1	0.196364	0.18
3	4	1/4/2011	1	0	1	0	2	1	1	0.200000	0.21
4	5	1/5/2011	1	0	1	0	3	1	1	0.226957	0.22
5	6	1/6/2011	1	0	1	0	4	1	1	0.204348	0.23
6	7	1/7/2011	1	0	1	0	5	1	2	0.196522	0.20
7	8	1/8/2011	1	0	1	0	6	0	2	0.165000	0.16
8	9	1/9/2011	1	0	1	0	0	0	1	0.138333	0.11
9	10	1/10/2011	1	0	1	0	1	1	1	0.150833	0.15
10	11	1/11/2011	1	0	1	0	2	1	2	0.169091	0.19
11	12	1/12/2011	1	0	1	0	3	1	1	0.172727	0.16
12	13	1/13/2011	1	0	1	0	4	1	1	0.165000	0.15
13	14	1/14/2011	1	0	1	0	5	1	1	0.160870	0.18
14	15	1/15/2011	1	0	1	0	6	0	2	0.233333	0.24
15	16	1/16/2011	1	0	1	0	0	0	1	0.231667	0.23
16	17	1/17/2011	1	0	1	1	1	0	2	0.175833	0.17
17	18	1/18/2011	1	0	1	0	2	1	2	0.216667	0.23
18	19	1/19/2011	1	0	1	0	3	1	2	0.292174	0.29
19	20	1/20/2011	1	0	1	0	4	1	2	0.261667	0.25
20	21	1/21/2011	1	0	1	0	5	1	1	0.177500	0.15
21	22	1/22/2011	1	0	1	0	6	0	1	0.059130	0.07
22	23	1/23/2011	1	0	1	0	0	0	1	0.096522	0.09
23	24	1/24/2011	1	0	1	0	1	1	1	0.097391	0.11
24	25	1/25/2011	1	0	1	0	2	1	2	0.223478	0.23
25	26	1/26/2011	1	0	1	0	3	1	3	0.217500	0.20
26	27	1/27/2011	1	0	1	0	4	1	1	0.195000	0.21
27	28	1/28/2011	1	0	1	0	5	1	2	0.203478	0.22
28	29	1/29/2011	1	0	1	0	6	0	1	0.196522	0.21
29	30	1/30/2011	1	0	1	0	0	0	1	0.216522	0.25
30	31	1/31/2011	1	0	1	0	1	1	2	0.180833	0.18
31	32	2/1/2011	1	0	2	0	2	1	2	0.192174	0.23

```
In [ ]: numeric_features = ['temp', 'atemp', 'hum', 'windspeed']
bike_data[numeric_features + ['rentals']].describe()
```

```
Out[ ]:
```

	temp	atemp	hum	windspeed	rentals
count	731.000000	731.000000	731.000000	731.000000	731.000000
mean	0.495385	0.474354	0.627894	0.190486	848.176471
std	0.183051	0.162961	0.142429	0.077498	686.622488
min	0.059130	0.079070	0.000000	0.022392	2.000000
25%	0.337083	0.337842	0.520000	0.134950	315.500000
50%	0.498333	0.486733	0.626667	0.180975	713.000000
75%	0.655417	0.608602	0.730209	0.233214	1096.000000
max	0.861667	0.840896	0.972500	0.507463	3410.000000

Esto garantiza que los gráficos se muestren en línea en el bloc de notas de Jupyter

Obtener la columna de etiqueta

Crear una figura para 2 subgráficos (2 filas, 1 columna)

Trazar el histograma

Agregar líneas para la media, la mediana y el modo

Trazar el diagrama de caja

Agregar un título a la figura

Mostrar la figura

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

%matplotlib inline

#
label = bike_data['rentals']

#
fig, ax = plt.subplots(2, 1, figsize = (9,12))

#
ax[0].hist(label, bins=100)
```

```
ax[0].set_ylabel('Frequency')

#
ax[0].axvline(label.mean(), color='magenta', linestyle='dashed', linewidth=2)
ax[0].axvline(label.median(), color='cyan', linestyle='dashed', linewidth=2)

#
ax[1].boxplot(label, vert=False)
ax[1].set_xlabel('Rentals')



#
fig.suptitle('Rental Distribution')

#
fig.show()
```

C:\Users\lenovo\AppData\Local\Temp\ipykernel_16684\2229708211.py:30: UserWarning: Matplotlib is currently using module://matplotlib_inline.backend_inline, which is a non-GUI backend, so cannot show the figure.


```
fig.show()
```


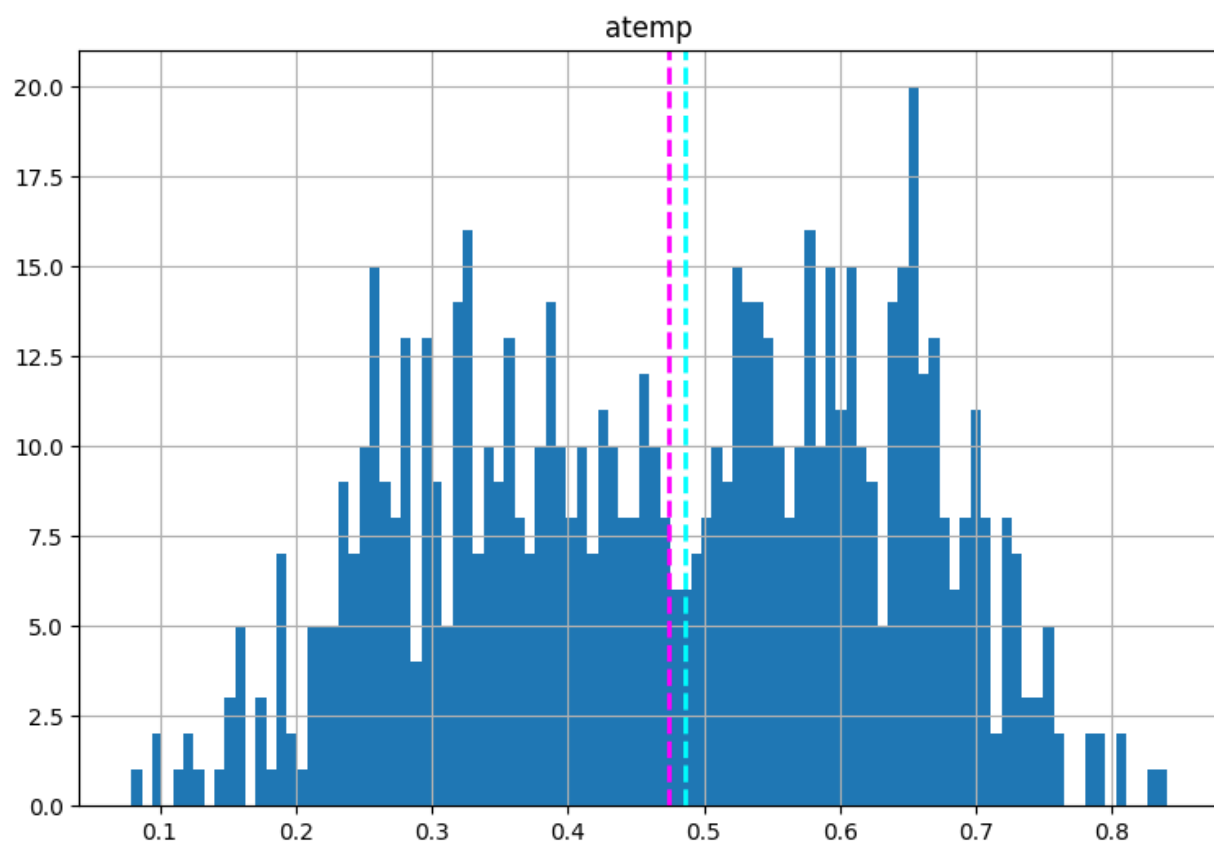
Rental Distribution




Trazar un histograma para cada entidad numérica

```
In [ ]: #
for col in numeric_features:
    fig = plt.figure(figsize=(9, 6))
    ax = fig.gca()
    feature = bike_data[col]
    feature.hist(bins=100, ax = ax)
    ax.axvline(feature.mean(), color='magenta', linestyle='dashed', linewidth=2)
    ax.axvline(feature.median(), color='cyan', linestyle='dashed', linewidth=2)
    ax.set_title(col)
plt.show()
```






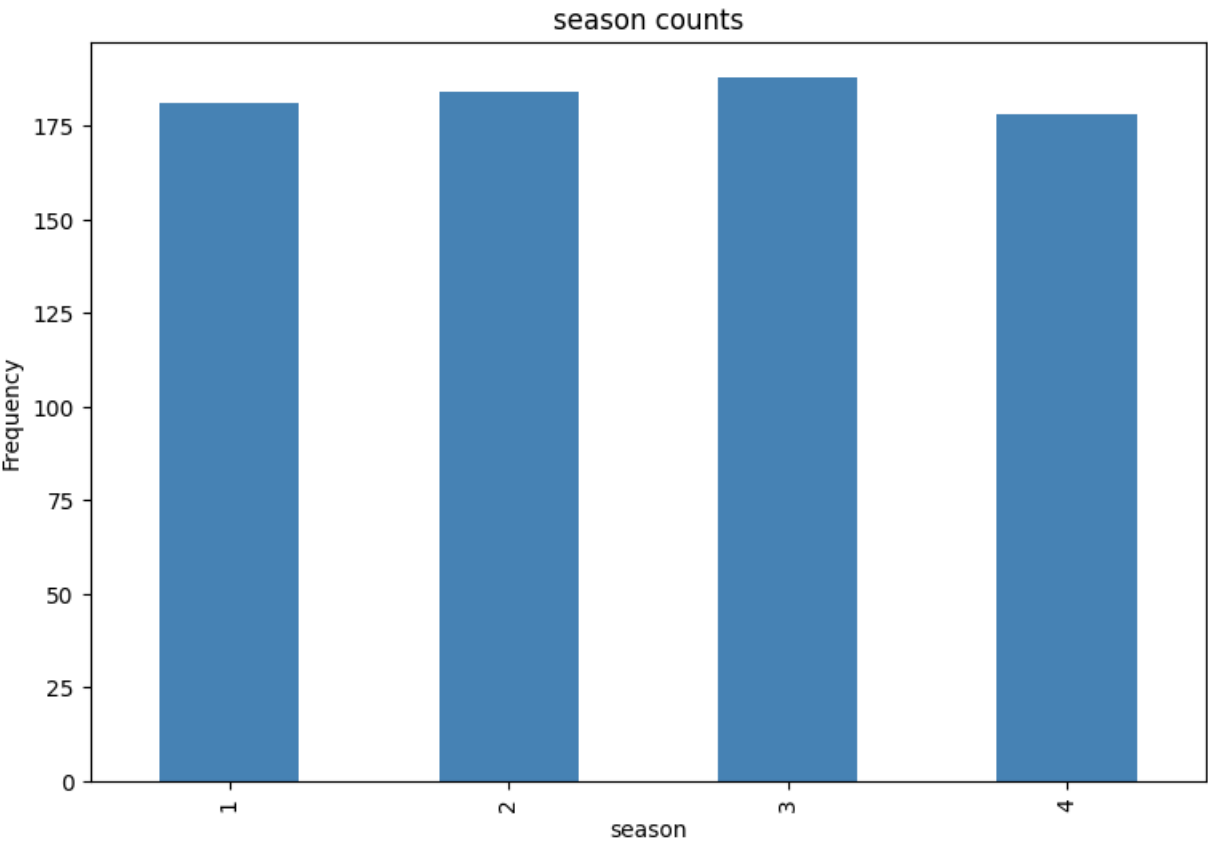


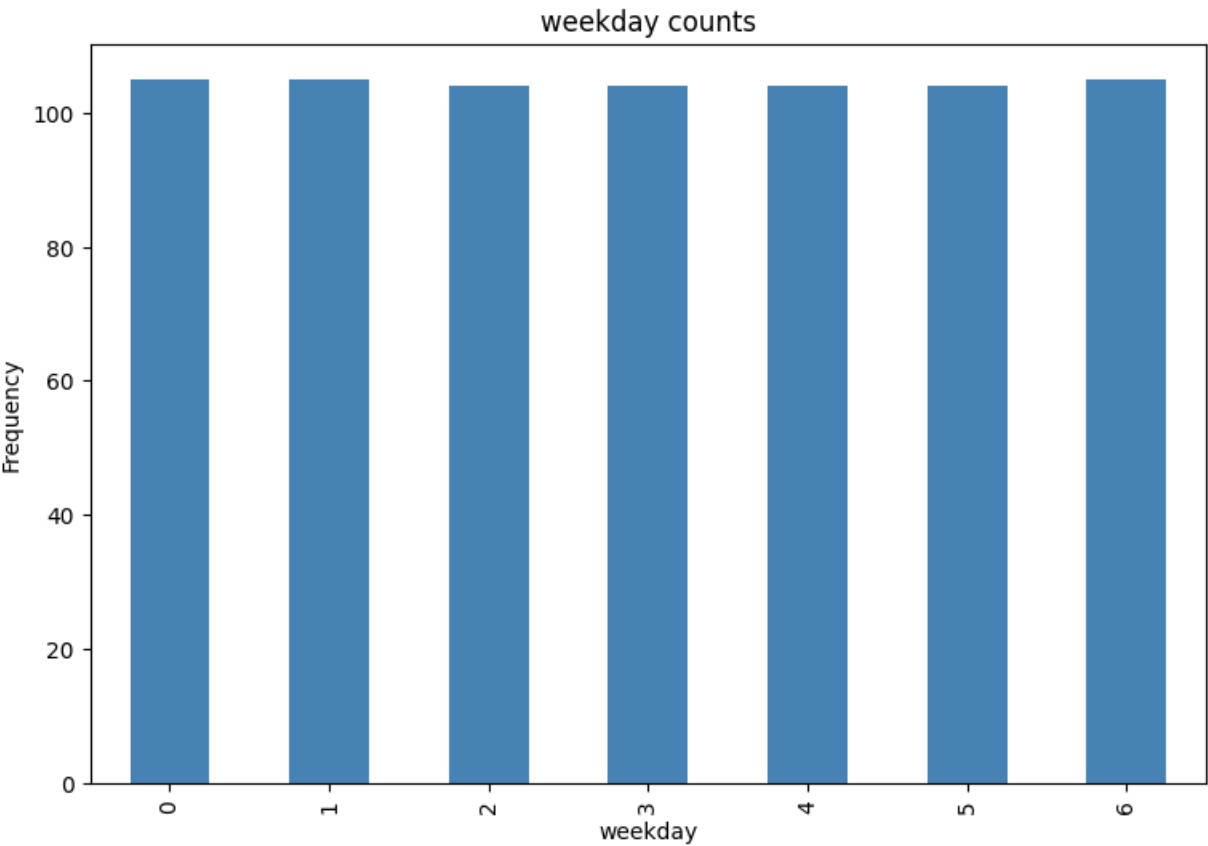

Trazar un gráfico de barras para cada recuento de características categóricas



```
In [ ]: import numpy as np


#
categorical_features = ['season', 'mnth', 'holiday', 'weekday', 'workingday', 'weathersit',

for col in categorical_features:
    counts = bike_data[col].value_counts().sort_index()
    fig = plt.figure(figsize=(9, 6))
    ax = fig.gca()
    counts.plot.bar(ax = ax, color='steelblue')
    ax.set_title(col + ' counts')
    ax.set_xlabel(col)
    ax.set_ylabel("Frequency")
plt.show()
```


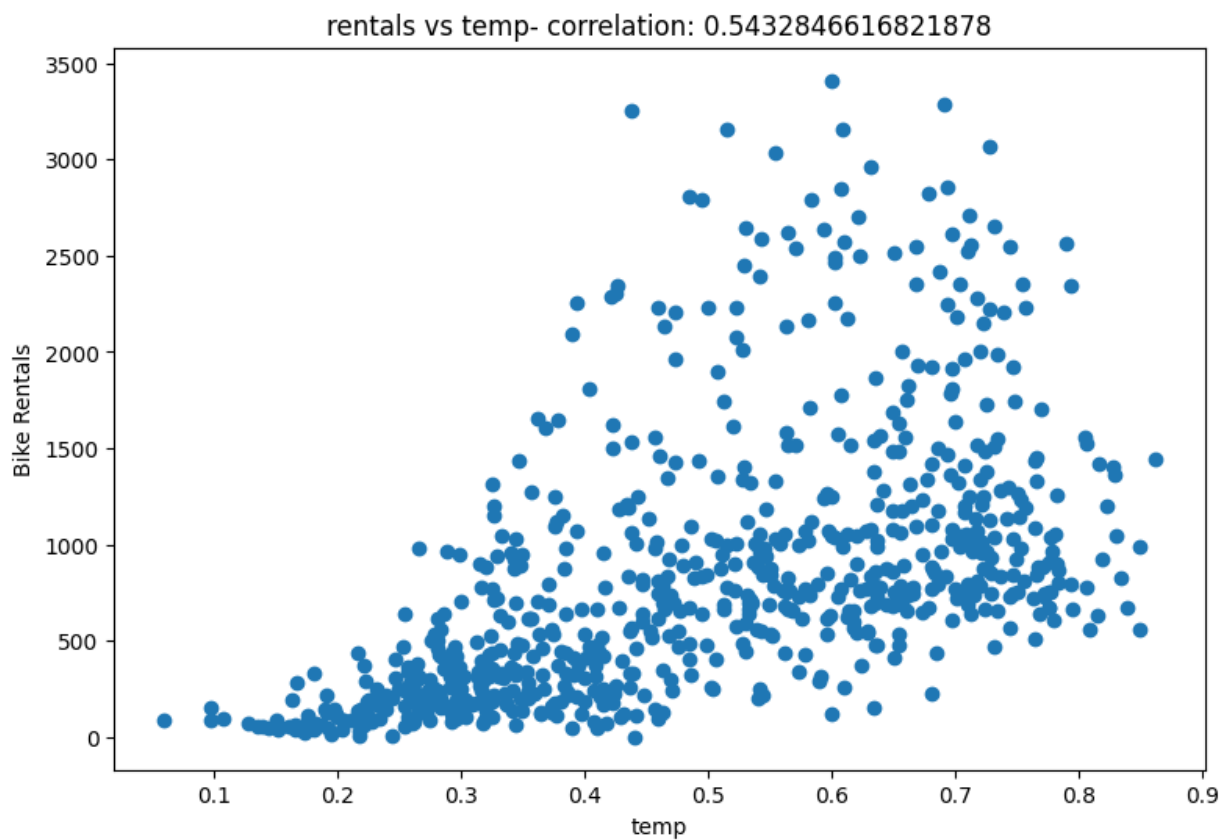




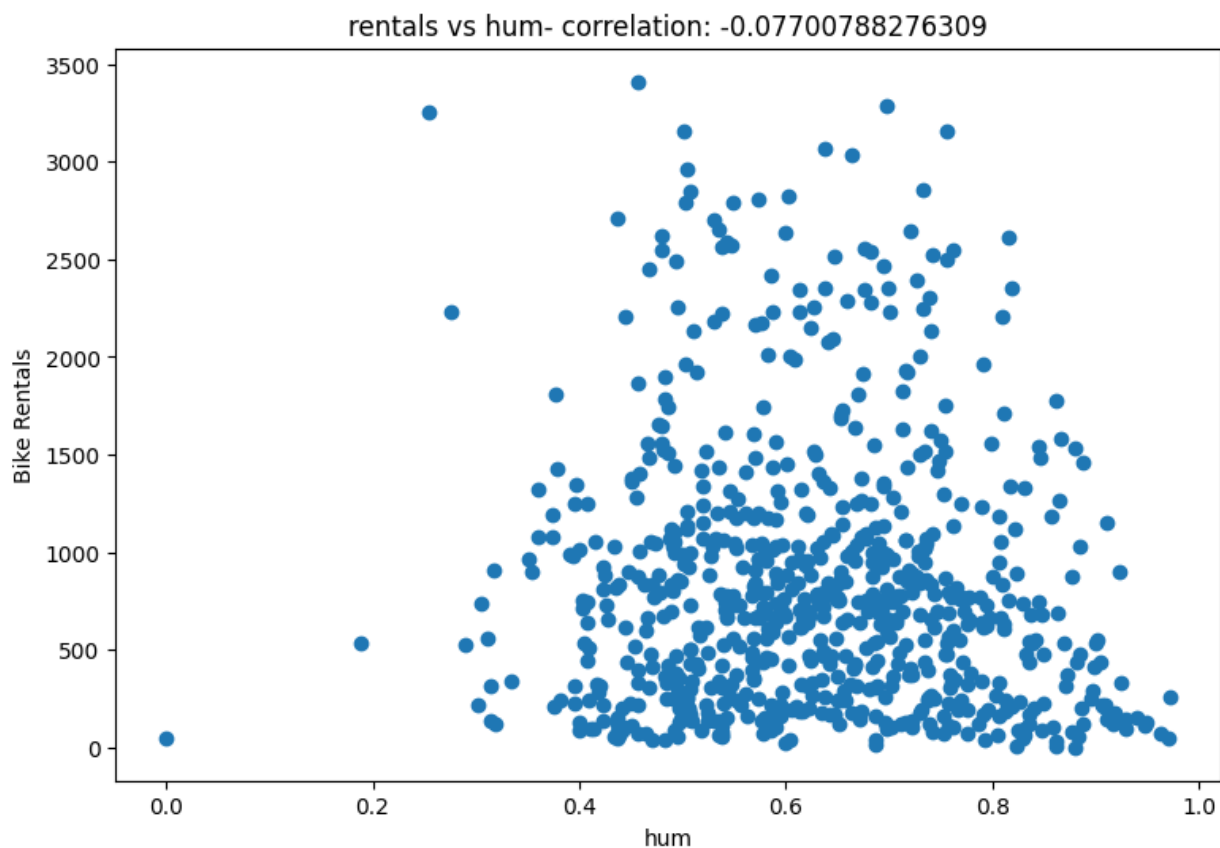







```
In [ ]: for col in numeric_features:
    fig = plt.figure(figsize=(9, 6))
    ax = fig.gca()
    feature = bike_data[col]
    label = bike_data['rentals']
    correlation = feature.corr(label)
    plt.scatter(x=feature, y=label)
    plt.xlabel(col)
    plt.ylabel('Bike Rentals')
    ax.set_title('rentals vs ' + col + '- correlation: ' + str(correlation))
plt.show()
```








Trazar un diagrama de caja para la etiqueta por cada característica categórica

```
In [ ]: #  
for col in categorical_features:  
    fig = plt.figure(figsize=(9, 6))  
    ax = fig.gca()  
    bike_data.boxplot(column = 'rentals', by = col, ax = ax)  
    ax.set_title('Label by ' + col)  
    ax.set_ylabel("Bike Rentals")  
plt.show()
```




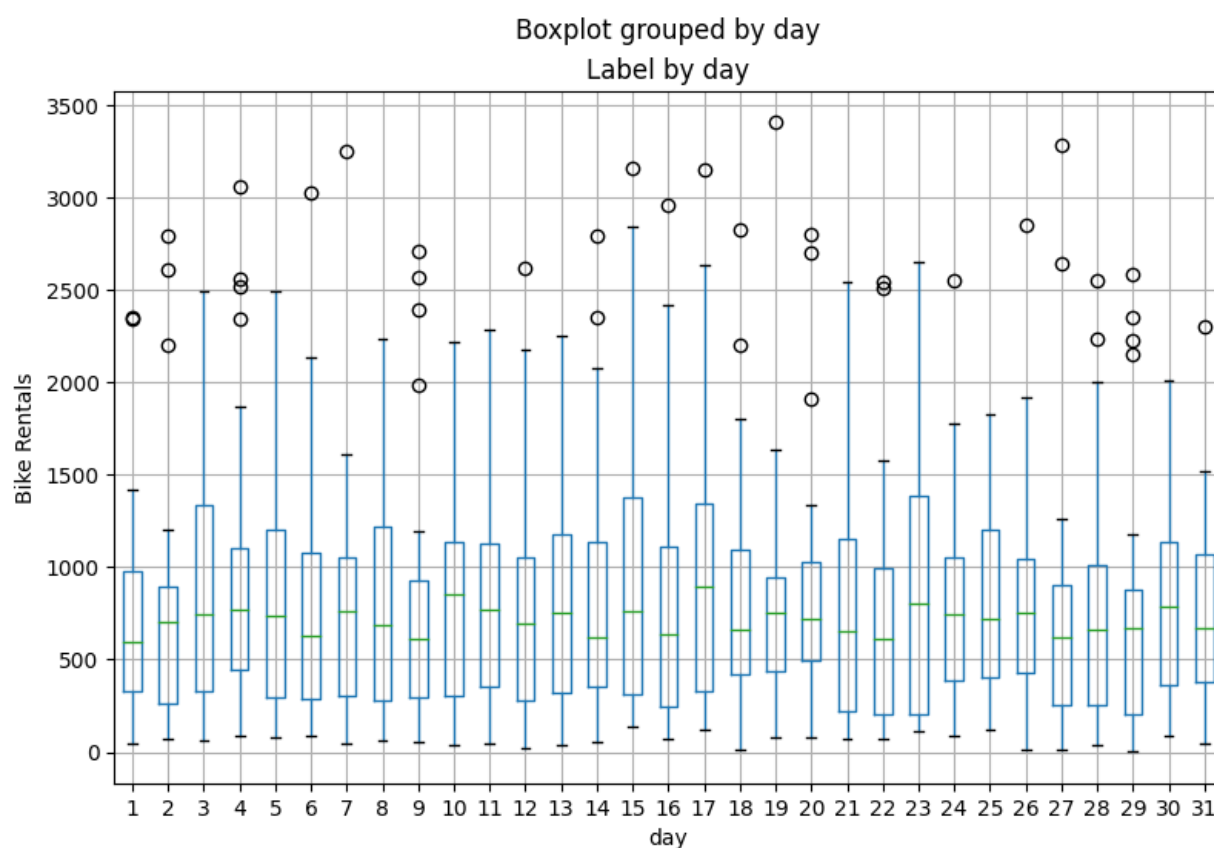



Boxplot grouped by weekday
Label by weekday



Boxplot grouped by workingday
Label by workingday





Características y etiquetas separadas

```
In [ ]: #
X, y = bike_data[['season', 'mnth', 'holiday', 'weekday', 'workingday', 'weathersit', 'temp
```

```
print('Features:',X[:10], '\nLabels:', y[:10], sep='\n')
```

Features:

```
[[1.      1.      0.      6.      0.      2.      0.344167
  0.363625 0.805833 0.160446 ]
 [1.      1.      0.      0.      0.      2.      0.363478
  0.353739 0.696087 0.248539 ]
 [1.      1.      0.      1.      1.      1.      0.196364
  0.189405 0.437273 0.248309 ]
 [1.      1.      0.      2.      1.      1.      0.2
  0.212122 0.590435 0.160296 ]
 [1.      1.      0.      3.      1.      1.      0.226957
  0.22927 0.436957 0.1869   ]
 [1.      1.      0.      4.      1.      1.      0.204348
  0.233209 0.518261 0.0895652]
 [1.      1.      0.      5.      1.      2.      0.196522
  0.208839 0.498696 0.168726 ]
 [1.      1.      0.      6.      0.      2.      0.165
  0.162254 0.535833 0.266804 ]
 [1.      1.      0.      0.      0.      1.      0.138333
  0.116175 0.434167 0.36195   ]
 [1.      1.      0.      1.      1.      1.      0.150833
  0.150888 0.482917 0.223267 ]]
```

Labels:

```
[331 131 120 108  82  88 148  68  54  41]
```

Divida los datos entre el 70% y el 30% en el conjunto de entrenamiento y el conjunto de prueba

```
In [ ]: from sklearn.model_selection import train_test_split

#
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)

print ('Training Set: %d rows\nTest Set: %d rows' % (X_train.shape[0], X_test.shape[0]))
```

Training Set: 511 rows

Test Set: 220 rows

Entrenar el modelo

Ajustar un modelo de regresión lineal en el conjunto de entrenamiento

```
In [ ]: #
from sklearn.linear_model import LinearRegression

#
model = LinearRegression().fit(X_train, y_train)
print (model)
```

LinearRegression()

```
In [ ]: import numpy as np
```

```

predictions = model.predict(X_test)
np.set_printoptions(suppress=True)
print('Predicted labels: ', np.round(predictions)[:10])
print('Actual labels   : ', y_test[:10])

```

```

Predicted labels: [1896. 1184. 1007. -28. 314. 385. 475. 590. 1476. -22.]
Actual labels   : [2418  754  222   47  244  145  240  555 3252   38]

```

Superponer la línea de regresión


```

In [ ]: import matplotlib.pyplot as plt

%matplotlib inline

plt.scatter(y_test, predictions)
plt.xlabel('Actual Labels')
plt.ylabel('Predicted Labels')
plt.title('Daily Bike Share Predictions')
#
z = np.polyfit(y_test, predictions, 1)
p = np.poly1d(z)
plt.plot(y_test, p(y_test), color='magenta')
plt.show()

```



```

In [ ]: from sklearn.metrics import mean_squared_error, r2_score

mse = mean_squared_error(y_test, predictions)
print("MSE:", mse)

```

```
rmse = np.sqrt(mse)
print("RMSE:", rmse)

r2 = r2_score(y_test, predictions)
print("R2:", r2)
```

MSE: 201972.55947035595

RMSE: 449.4135728595165

R2: 0.604045473691919